a16z.com /2018/07/20/after-product-market-fit-marc-andreessen-elad-gil/

Where to Go After Product-Market Fit: An Interview with Marc Andreessen

22-28 minutes

Editor's note: This interview with Marc Andreessen was edited and condensed for clarity from the original conversation, and appears in The High Growth Handbook on scaling companies from 10 to 10,000 by Elad Gil. The full content has been reprinted below, with our own formatting. You can also listen to an a16z Podcast with Gil on the book (and other topics) here.

It's a thrilling thing to build a new product, then watch as consumers actually pick it up. But achieving product/market fit also marks the beginning of a challenging time for a lot of founders. You've poured your energy into getting here — now what?

Few people have as much insight into this make-or-break moment as Marc Andreessen. As a repeat founder himself and one of Silicon Valley's most influential investors, he's seen firsthand that the decisions startups make at this juncture are some of the most consequential they will ever make. I spoke with Marc to get his top recommendations for how startup leaders can turn early success into lasting relevance.

Elad: After you've achieved product/market fit, what do you think are the most important determinants of a company's success? You have your first product working, everything is scaling, everything seems to be going great, but now it's time for you to do these three things. What are those things in your mind, or what are the most common issues people run into?

Marc: I think there are three big categories. Once there's product/market fit, then the main thing becomes taking the market — which is to say, figuring out how to get the product to the entire market, how to get dominant market share; because most tech markets tend to end up with one company with most of the market share. And that company tends to be all the value that gets created in that sector, from a return standpoint. That company also tends to have all the resources to do everything else that they want to do, including build new products.

So winning the market is the big thing. The thing that is so essential that people need to understand is that the world is a really big place. The good news is that markets are bigger than ever. There are more consumers on the internet than ever before. There are more businesses that use software than ever before.

There are more sectors of the economy where this stuff all matters. And so the markets are bigger than before. But that means that the challenge of building an organization, a model, and a distribution capability that can actually get the product to all the customers is an intense challenge. And of course the personality type of the technical founder who creates a breakthrough product, they don't necessarily intuitively understand that that next part involves taking down the market.

That's number one. Number two is getting to the next product. We are in a product cycle business. Which is to say that every product in tech becomes obsolete, and they become obsolete pretty quickly. If all you do is take your current product to market and win the market, and you don't do anything else — if you don't keep innovating — your product will go stale. And somebody will come out with a better product and displace you.

So you do need to get to the next product. Of course that's a punishingly hard thing to do. It was hard enough to get to the first one, and to come up with the second one is often even harder. Although the consistency between these two tasks is this: If you do take the market, you tend to have the financial resources to be able to invest heavily in R&D. And you also develop M&A currency, so you can then go buy the second product if you have to. It gives you another option to get to the second product.

Elad: And that can also use your established distribution that you can just plug into, which is nice.

Marc: Exactly. In fact, the general model for successful tech companies, contrary to myth and legend, is that they become distribution-centric rather than product-centric. They become a distribution channel, so they can get to the world. And then they put many new products through that distribution channel. One of the things that's most frustrating for a startup is that it will sometimes have a better product but get beaten by a company that has a better distribution channel. In the history of the tech industry, that's actually been a more common pattern. That has led to the rise of these giant companies over the last fifty, sixty, seventy years, like IBM, Microsoft, Cisco, and many others.

> The general model for successful tech companies, contrary to myth and legend, is that they become distribution-centric rather than product-centric. They become a distribution channel, so they can get to the world.

But then the third thing you need to do is what I call "everything else," which is building the company around the product and the distribution engine. That means becoming competent at finance, HR, legal, marketing, PR, investor relations, and recruiting.

That's the stuff that's the easiest to put to one side — for a little while. If you've got a killer product and a great sales engine, you can put that other stuff aside for a while. But the longer you put that stuff aside, the more risk that you develop and the more you expose yourself to catastrophic failure through self-inflicted wounds.

Of course, the obvious one that we're all seeing right now in this environment is HR. The number of companies in the Valley that put HR off to the side and decided it wasn't important and are now dealing with some level of catastrophe — either a public catastrophe or one that's in the making — that's a pretty high percentage of companies right now. And it's totally unnecessary. If they had taken HR seriously starting at an earlier point, they probably would have been able to fight a lot of their issues. But for whatever reason they decided it wasn't important. So HR has to be taken seriously.

Then there's legal. We see cases in the Valley where companies have just decided that laws are optional, felonies are fine. And at some point, not having a general counsel who's able to explain to the CEO where the line is, that becomes a big issue.

And obviously finance. There are companies that blow themselves up financially that don't have to, with wildly out-of-control cost structures or horribly screwed-up pricing or whatever.

Elad: A common question for a lot of founders, especially first-time founders, is when should they actually hire an HR leader or somebody in the HR function or a GC [general counsel] or somebody for finance? Is that at a certain round of financing? Is it at a certain revenue level? Is it at a certain number of people? When should you start adding those different functions?

Marc: It's somewhere between 50 to 150 people. It's somewhere in there. If you don't start layering in HR once you've passed 50 people on your way to 150, something is going to go badly wrong.

I think there's actually an explanation for that. It's because 150 is the Dunbar number, the number of people you can directly know. So somewhere between 50 to 150 people, everybody doesn't know everybody. There are people running around who other people have never met. When you were 5 or 10 or 20 people, it was one big happy family and everybody knew everybody — well, maybe it wasn't, but at least everybody knew everybody. And the CEO had direct one-on-one relationships with everybody in the company. Once you get to 50+, that's just no longer the case. At that point there's a necessary impersonality of the professional relationships in the company. And then HR catastrophes emerge, because people trip the line from proper professional behavior in a workplace to improper. There are just too many people running around interacting.

Elad: Let's focus on each of these three points. When it comes to taking down the market, can you talk a little bit more about the primary things that people should be doing? And what do people tend to miss, or where do they tend to screw things up?

Marc: I think the single biggest thing is that every market has early adopters. There are early adopters for everything, and it's kind of amazing that that's the case. But there are always people. There are people on product sites every day looking for the next new consumer thing to try. There are even early adopter CIOs. There are Fortune 500 CIOs who pride themselves on discovering the next new whatever it is — the new relational database, AI, whatever it is. They're all over it.

And a lot of product/market fit is the fit with the early adopters. And so you get these extremely enthusiastic people, who in a lot of cases have sought you out as the vendor, saying, "Wow, your thing is really cool. Can I please use it?" And that's your sign of product/market fit.

The problem is, the early adopters are only ever a small percentage of the overall market. And so a lot of founders, especially technical ones, will convince themselves that the rest of the market behaves like the early adopters, which is to say that the customers will find them. And that's just not true.

In the consumer world, it's not true because people have plenty of existing things they can spend their time on. They have to be convinced to try the next new thing. And so whether you want to call that marketing or growth hacking or user acquisition or whatever you want to call it, there's some distribution function there, for all these things, that's critical.

And then for sure for B2B. Most businesspeople in the world, most CIOs or whoever is going to buy technology in a business, don't wake up in the morning and say, "Gee whiz, can I go find the next hot thing to take a chance on?" That's not how businesspeople live their lives. And so there needs to be some distribution capability to get to them.

Again, the danger there is a market share thing. If you stick to the early adopters, you'll get 5% of the market, but you're not going to get 95% of the market. And that means, sort of by definition, somebody else is going to go get 95%.

One of the things you see crystal clearly in VC is how much competition emerges whenever anything works. Every single time we say, "Oh, this startup is unique. There's some unique product here and there's not going to be competition," invariably six months later there are 20 venture-backed competitors doing the exact same thing. And so at some point, if the early guys don't get to the other 95% of the market, somebody else is going to go take it away. And whoever has 95% of the market, number one they're going to get all the value. All the investment returns, all the employee compensation flows to that company. And then number two, that company then accretes resources so they can work backward. In a lot of cases, they end up buying the company that got the early adopters for a small percentage of their equity, and then they just take the whole thing.

Elad: You mentioned three or four tactical things that startups can do to stay viable. One is product iteration and building products that serve more of that market. Second, you really emphasize building up distribution. Third is M&A, which seems to be really underutilized in Silicon Valley today, at least in terms of the next generation of companies. If you have a $10 or $20 billion market cap, you should be buying things. One thing you didn't mention, but I'd love to hear more on, is moats — in other words, building defensibility into what you're doing. As you think about those four different factors, how do you rank them or think about common failure modes?

Marc: I am shocked by the absence of M&A relative to what I would expect in the environment. And I would say there's no question that the big new tech incumbents are not buying enough stuff just on the math. I think it's just kind of obvious. In the old days, their predecessor companies were far more aggressive at building up their positions for M&A. And honestly, the Fortune 500, the big public companies, are not nearly as aggressive as I think they're going to be

I think this is a temporary lull. I think five years from now we're going to be having a very different conversation, because it's just going to become obvious that this is an underutilized thing. And I do think that that could be a very effective weapon, therefore, for somebody who really figures this out and does it aggressively in the right way.

Cisco is one of the great case studies in the Valley. It's a very successful, very big, very established company, and a very large percentage of that has been M&A. And then obviously Google. Probably an under-told part of the Google story is how M&A built Google. People, I think, don't even necessarily remember the number of things that Google bought that turned into what you think of today as Google-originated products.

As far as defensibility, I think you construct defensibility through some combination of product innovation and distribution building. You construct it. You obviously want as much defensibility as you can get in your product, and so you try to get as far out ahead as you can. It's the idealized Peter Thiel model of "build something nobody else can build." Or the SpaceX model of "go get all the talent."

The problem with that is true defensibility purely at the product level is really rare in the Valley, because there are a lot of really good engineers. And there are new ones every day, whether they're coming out of Stanford or coming in from other countries or whatever. And then there's the issue of leap-frogging. The next team has the opportunity to learn from what you did and then build something better. So I think pure product defensibility is obviously highly desirable, but it's actually quite difficult.

I think the distribution moats end up being at least as important. At some point, whoever has the distribution engine and gets 100% of the market, at some point that engine itself is a moat. Again, that

might be an enterprise sales team for a SaaS company, or it might be the growth team at a consumer company.

One interesting question I have is: Would you rather have another two years' lead on product, or a two years' lead on having a state-of-the-art growth effort? I think the answer for a lot of consumer products is actually that you'd rather have the growth effort.

The other big missing variable in all of this is pricing. I've talked in public about this before. What I don't hear from companies is, "Oh, we don't think we have a moat." What I hear from companies is, "Oh, we have an awesome moat, and we're still going to price our product cheap, because we think that's somehow going to maximize our business." I'm always urging founders to raise prices, raise prices, raise prices.

First of all, raising prices is a great way to flesh out whether you actually do have a moat. If you do have a moat, the customers will still buy, because they have to. The definition of a moat is the ability to charge more. And so number one, it's just a good way to flesh out that topic and really expose it to sunlight.

> Raising prices is a great way to flesh out whether you actually do have a moat… The definition of a moat is the ability to charge more.

And then number two, companies that charge more can better fund both their distribution efforts and their ongoing R&D efforts. Charging more is a key lever to be able to grow. And the companies that charge more therefore tend to grow faster.

That's counterintuitive to a lot of engineers. A lot of engineers think there's a one-dimensional relationship between price and value. They have this mental model of commerce like they're selling rice or something. It's like, "My product is magical and nobody can replicate it, and I need to price it like it's a commodity." No, you don't. In fact, quite the opposite. If you price it high, then you can fund a much more expensive sales and marketing effort, which means you're much more likely to win the market, which means you're much more likely to be able afford to do all the R&D and acquisitions you're going to want to do. And so we always try to snap people into a two-dimensional mindset, where higher prices equals faster growth.

Elad: That's an awesome insight. I feel like there's two really key notes that you brought up that typically aren't talked about. One is distribution moats. I think people emphasize network effects and data effects way too much, and I've never seen a real data effect, at least recently. And then second is charging more equals faster growth. Those are really key things that people really don't talk about or think about.

Marc: I think network effects are great, but in a sense they're a little overrated. The problem with network effects is they unwind just as fast. And so they're great while they last, but when they reverse, they reverse viciously. Go ask the MySpace guys how their network effect is going. Network effects can create a very strong position, for obvious reasons. But in another sense, it's a very weak position to be in. Because if it cracks, you just unravel. I always worry when a company thinks the answer is just network effects. How durable are they?

> The problem with network effects is they unwind just as fast. And so they're great while they last, but when they reverse, they reverse viciously.

To your point on data network effects, I would just say that we don't see it very often. We see a lot of claims, and very little evidence. The reality is, there's a lot of data in the world, and a lot of ways to get data. We have not seen very many data moats that actually make sense, even in science. Deep learning is the latest area where people think there's data network effects. The problem is there's innovation in deep learning to actually do deep learning on small data sets now. So even the science strains underneath that in a way that's undermining it. So that's risky.

Elad: I'd love to talk a little bit about getting the next product in the product cycle. How do you start iterating and how do you come up with your v2 or your new product area? And how do you think about percent investment in core adjacent versus completely new areas? Google had a 70-20-10 framework. Do you think frameworks like that work?

Marc: I don't really like the numeric version of the answer because it's kind of what big, dumb companies do. They say, well, we invest R&D as a percentage. But anybody who's actually worked in R&D knows it's not really a question of money. It's not really a question of percentage of spend. It's who's doing it.

What I've always found is this: give me a great product picker and a great architect, and I'll give you a great product. But if I don't have a great product manager, a great product originator — it used to be

called a product picker — and I don't have a great architect, I'm not going to get a great product.

Elad: Google's framework, by the way, was percentage of people. So it was 70% of human resources versus financial. But fair enough.

Marc: But it's kind of the same thing. I mean, it's a fine concept, but it begs the question: Who are the people?

I would take more of a micro-view of it. Which is: Okay, how many great product pickers do you have, people who can actually conceptualize new products? And then how many great architects do you have, who can actually build it? Sometimes, by the way, those are the same person. Sometimes it's a solo act. And sometimes that's the founder.

As you scale, you need more of those people. But I always think it's a matter of, okay, how many of those people do you have or can you go get? Or, back to acquisition, how many of those can you acquire? And then basically that's the number of products you can be working on. You organize R&D around that, in my view. You want to have a relatively flat R&D structure. You basically want to have autonomous teams, where each team is guaranteed to have a great product person and a great architect. And that's the model.

This is why I always ask people, okay, let's just do an inventory of how many of those people you think you have. And even at really, really big companies, it's not a large number. Even at giant companies there might be ten or twenty of each, maybe. And then you build the rest of the engineering organization around those people, including all the rest of the stuff that you do for recruiting and onboarding and all that other stuff. But at the core of it, who do you actually have who can conceptualize new products, and who do you actually have who can build it?

If it's the founder, fair enough. But you need to construct the organization so that the founder has the time to continue to do that. So that gets to all the questions around when you need an outside CEO or when you need a COO. And then, even the founders who can do that themselves eventually run out of time. And the challenges get bigger. So how do they attract and retain people who can pick up some of that work for them?

Elad: If it does get distributed, do you have any perspective on whether it should be a general management, verticalized structure versus a matrixed one?

Marc: I generally think matrixed is death, so I'm always pushing companies to go to a flat structure of independent teams. I'm really on the Jeff Bezos program on that, the two-pizza team thing [Jeff Bezos favors product innovation teams with five to seven people — no more than can be fed with two pizzas]. I think hierarchies kill innovation for the most part. And I think that matrixes are just lethal in most cases. There are exceptions, but in most cases, you need original thinking and speed of execution, and it's really hard to get that in anything other than a small-team format, in my view.